

# Lecture 18

---

COP3402 FALL 2015 – DR. MATTHEW GERBER – 11/23/2015

FROM EURIPIDES MONTAGNE, FALL 2014


# Tonight

---

- Review: The Tiny Computer
- The Interrupt Cycle
- Memory Protection
- Privileged Instructions
- System Calls

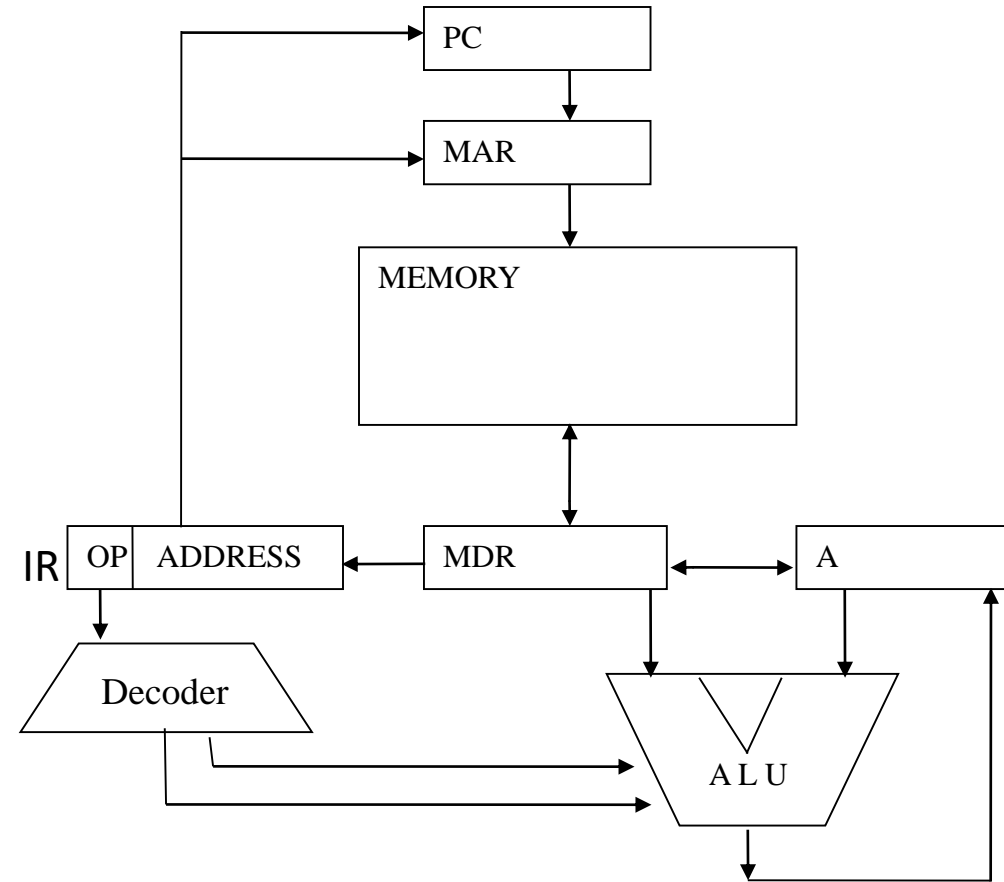
# Definitions, Part 1: Components

- **Main Storage (MEMORY):** Used to store programs and data; in modern systems this is invariably implemented as RAM
- **Arithmetic Logic Unit (ALU):** Performs the actual execution of instructions
- **Decoder:** Examines the IR (next slide), decides which instruction the processor will execute and signals the ALU to execute it



# Definitions, Part 2: Registers

- **Program Counter (PC):** Holds the address of the next instruction to be executed
- **Memory Address Register (MAR):** Used to specify the address to a specific memory location in main storage
- **Memory Data Register (MDR):** Used to store data being sent to or received from main storage
- **Instruction Register (IR):** Stores the instruction to be executed by the processor
  - Has two components that we'll discuss later
- **Accumulator (A):** Stores data to be used as input to the ALU



# The Instruction Cycle

---

The *instruction cycle*, *machine cycle* or *fetch-execute cycle* of the Von Neumann machine has two looping steps:

- **Fetch Cycle:** Retrieve an instruction from memory
- **Execute Cycle:** Execute the retrieved instruction

When the instruction is finished executing, the next instruction is fetched and the cycle continues.

# The Tiny ISA

0 Fetch (hidden instruction) 1 LOAD

MAR  $\leftarrow$  PC

MDR  $\leftarrow$  MEM[MAR]

IR  $\leftarrow$  MDR

PC  $\leftarrow$  PC+1

DECODER  $\leftarrow$  IR.OP

2 ADD

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  MEM[MAR]

A  $\leftarrow$  A + MDR

DECODER  $\leftarrow$  00

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  MEM[MAR]

A  $\leftarrow$  MDR

DECODER  $\leftarrow$  00

3 STORE

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  A

MEM[MAR]  $\leftarrow$  MDR

DECODER  $\leftarrow$  00

7 HALT

Op	Mnemonic	Meaning
0	(Fetch cycle)	Fetch next instruction
1	LOAD <addr>	Load <addr> contents into A
2	ADD <addr>	Add <addr> contents into A
3	STORE <addr>	Store A contents into <addr>
4	SUB <addr>	Sub <addr> contents from A
5	IN <device>	Read into A from <device>
6	OUT <device>	Write to <device> from A
7	HALT	End program
8	JMP <addr>	Branch to <addr>
9	SKIPZ	Skip next instruction if Z (A = 0)
A	SKIPG	Skip next instruction if G (A > 0)
B	SKIPL	Skip next instruction if L (A < 0)

# Error Detection

So far, we have a computer that can execute programs – if they're error-free.

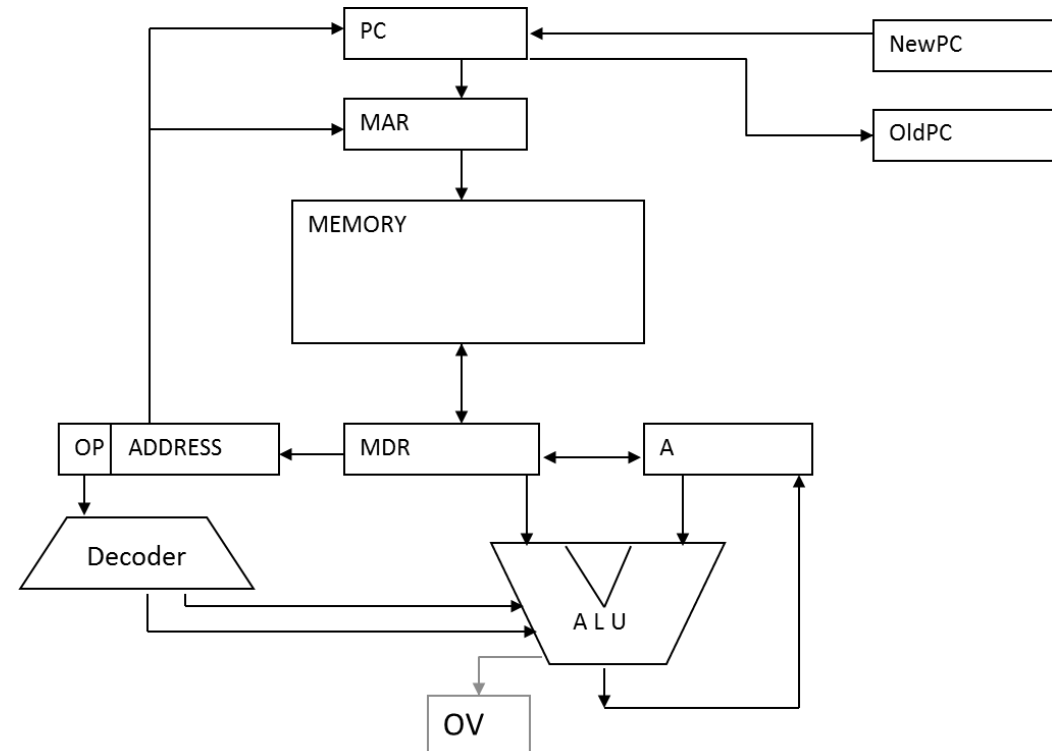
- What would happen if an overflow occurred while executing an addition operation?
- Let's think about a way to detect this type of event and act accordingly...

We can easily add a hardware flip/flop to the ALU that detects overflow.

- So far, so good
- But what do we *do* about it?

We need to detect this condition, and when it occurs, end the program abnormally.

- The Fetch/Execute cycle becomes a *Fetch/Execute/Interrupt* cycle.



# The Interrupt Cycle

---

0 Fetch (hidden instruction) 1 LOAD

MAR  $\leftarrow$  PC

MDR  $\leftarrow$  MEM[MAR]

IR  $\leftarrow$  MDR

PC  $\leftarrow$  PC+1

DECODER  $\leftarrow$  IR.OP

2 ADD

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  MEM[MAR]

A  $\leftarrow$  A + MDR

DECODER  $\leftarrow$  0F

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  MEM[MAR]

A  $\leftarrow$  MDR

DECODER  $\leftarrow$  0F

3 STORE

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  A

MEM[MAR]  $\leftarrow$  MDR

DECODER  $\leftarrow$  0F

7 HALT

Add a second hidden instruction (opcode F).

- Now, at the end of each execution cycle, the decoder is changed to F instead of 0

Instruction F is pretty simple:

F Abend (hidden instruction)

If OV=1

then HALT

DECODER  $\leftarrow$  00



# Interrupt Handling

Instead of simply halting execution we can transfer to an *interrupt handler routine*.

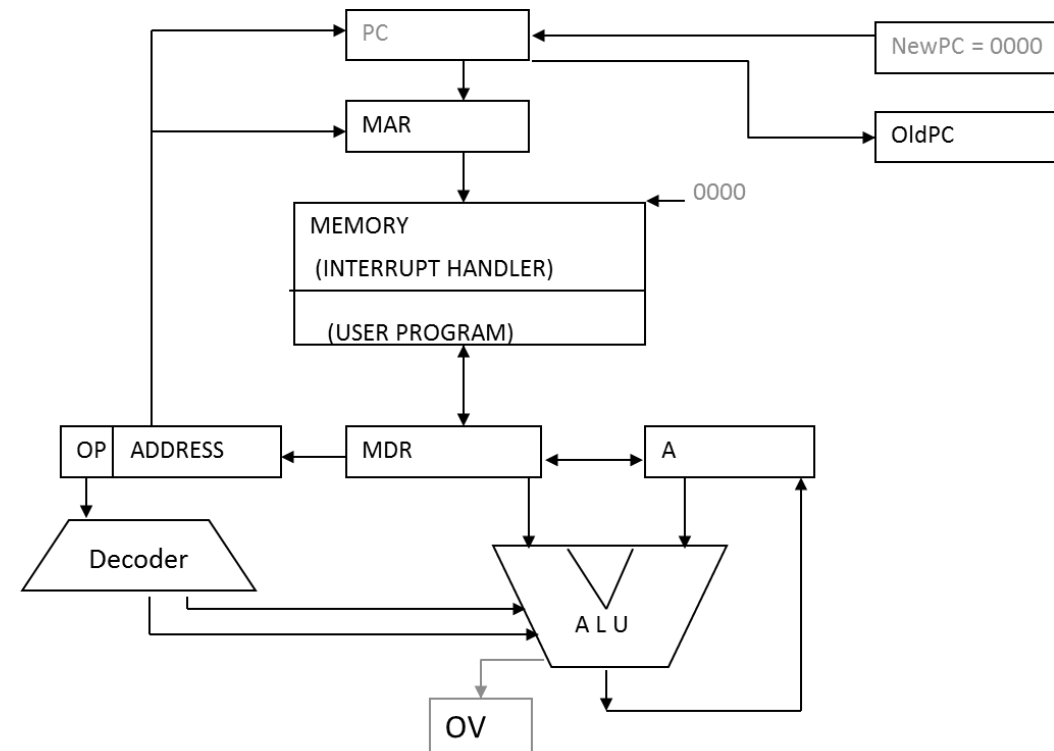
To do this, instead of halting, we just load the PC with the start address of the interrupt handler. For now, we'll call that **NewPC**.

F Handler (hidden instruction)

If  $OV=1$

then  $PC \leftarrow \text{NewPC}$

$\text{DECODER} \leftarrow 00$



# Interrupt Handlers: Implications

---

Along with LOAD, STORE, the I/O instructions, and the fetch cycle, the interrupt handler serves as part of a bridge between hardware and software.

- At least as importantly, with a little more work, it can serve as a bridge between *user* software and *system* software
- *The interrupt handler is one of the first steps toward a modern operating system*
- But before we can do much with it, we'd better put some walls around it


# Shared Memory

---

Nothing executes unless it's in memory, and that includes the interrupt handler.

So let's take a simple example.

- The interrupt handler occupies spaces 3000-3999
- The user program begins at space 4000
- What happens if the user writes to space 3500?



# Shared Memory

Nothing executes u  
that includes the i

So let's take a simp

- The interrupt han  
3999
- The user program
- What happens if

**BAD  
THINGS**

3500

4000

# Shared Memory

---


Nothing executes unless it's in memory, and that includes the interrupt handler.

So let's take a simple example.

- The interrupt handler occupies spaces 3000-3999
- The user program begins at space 4000
- What happens if the user writes to space 3500?

We need a way to protect the interrupt handler routine from user programs stepping on it.

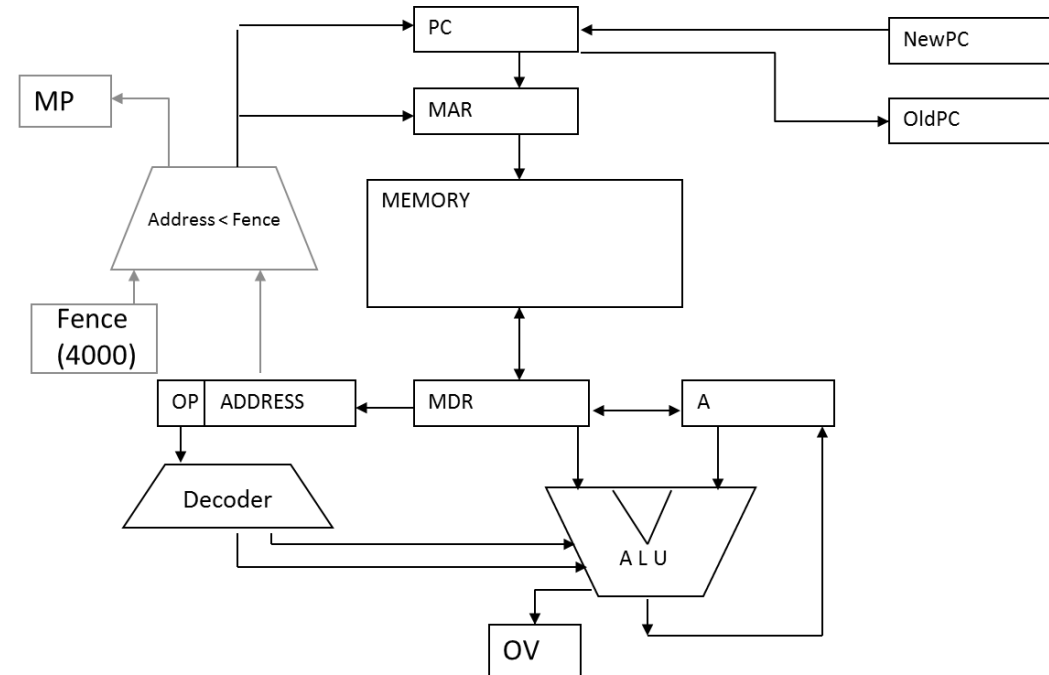
- In fact, it'd be great if we could protect *all* programs from *all other* programs stepping on them
- But let's take this a step at a time



# Simple Memory Protection

Our memory protection scheme has three components:

- A **Fence Register** loaded with the address of the boundary between the interrupt handler routine and the user program
- An **Address Comparator** that compares the fence register with any addresses that the user program attempts to access
- A **flip-flop** that is set to 1 if a user program tries to access something it shouldn't



# Instruction Changes

---

0 Fetch (hidden instruction) 1 LOAD

MAR  $\leftarrow$  PC

MDR  $\leftarrow$  MEM[MAR]

IR  $\leftarrow$  MDR

PC  $\leftarrow$  PC+1

DECODER  $\leftarrow$  IR.OP

MAR  $\leftarrow$  IR.Address

**IF MP=0 then**

MDR  $\leftarrow$  MEM[MAR]

A  $\leftarrow$  MDR

DECODER  $\leftarrow$  0F

2 ADD

MAR  $\leftarrow$  IR.Address

MDR  $\leftarrow$  MEM[MAR]

A  $\leftarrow$  A + MDR

DECODER  $\leftarrow$  0F

3 STORE

MAR  $\leftarrow$  IR.Address

**IF MP=0 then**

MDR  $\leftarrow$  A

MEM[MAR]  $\leftarrow$  MDR

DECODER  $\leftarrow$  0F

7 HALT

Now that we've got memory protection hardware, we need to use it.

- LOAD and STORE need to honor it
- The handler needs to react to a violation
- We treat a violation as an interrupt

F Handler (hidden instruction)

If OV=1 then PC  $\leftarrow$  NewPC

If MP=1 then PC  $\leftarrow$  NewPC

DECODER  $\leftarrow$  00

# Still Not Done

---

- ...but *what if a user program modifies the fence register?*
- (Yes, programmers will, in fact, do that if we let them.)
  - (Really.)



# Privileged Instructions & System Modes

---

We've protected memory, and that's all well and good... but *what would happen if a user program tried to modify the memory fence register?*

- The register isn't in memory – memory protection doesn't help it
- We need the idea of *privileged instructions* – instructions that user programs cannot access

To distinguish between times when privileged instructions are allowed and not, the computer operates in two **modes**.

- User mode: 0
- Supervisor mode: 1

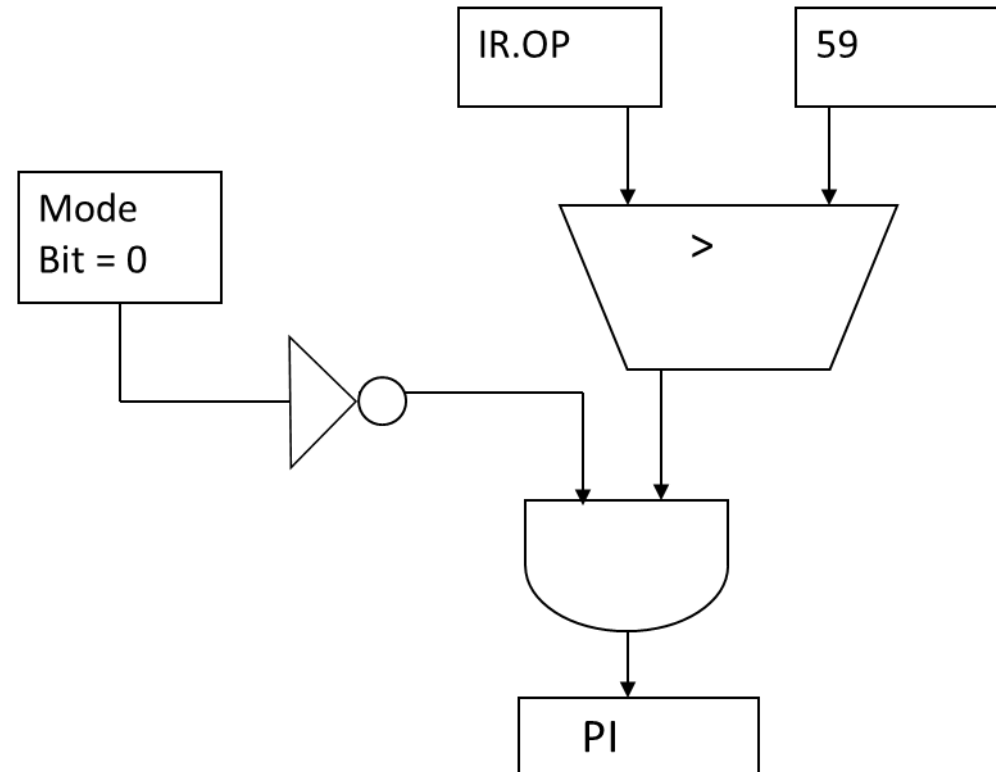
Supervisor mode is a mode exclusively reserved for the interrupt handler.

- In fact, *supervisor* is another name for the interrupt handler
- In user mode, only a subset of the instruction set can be used
- The supervisor has access to all instructions

# Privileged Instructions: Implementation

We have to add a few more pieces to the CPU.

- First, another flip-flop: the *mode bit*
- Second, another comparator
  - We organize the instruction set so that all and only privileged instructions have opcodes higher than some  $n$
  - Here  $n = 59$
  - The comparator compares the opcode with  $n$
- If the comparator finds that  $IR.OP > n$ , then the instruction is privileged
- If the instruction is privileged *and* the mode bit is 0 – indicating user mode – then we have an invalid instruction
  - And – you guessed it – we call the interrupt handler



# Types of Interrupts

---

There are several types of interrupts. So far we've been referring to ones created by errors – such as an overflow, memory address violation, or invalid instruction.

- Whenever one of these occurs, the supervisor *will* abort the process

Obviously, that's not the only type of interrupt.

- **I/O interrupts** occur when a device sends a signal to inform the CPU that an I/O operation has been completed
- **External** interrupts occur when non-I/O system components – such as the timer – send signals to the CPU
- **System calls** are interrupts generated *intentionally* by user processes, to request services from the operating system

When interrupts other than errors occur, the state of the running process must be saved before the supervisor services them.

# Program State Word (PSW)

---

The PSW **Program State Word** is a structure in the CPU that holds information about the state of a program.

In this structure, we have (at least):

- The program counter
- The interrupt flags we've discussed
  - OV (Overflow)
  - MP (Memory Protection)
  - PI (Privileged Instruction)
  - I/O
- The mask, which we'll define later
- The system mode

PC	Interrupt Flags						MASK	Mode
	OV	MP	PI		I/O		To be defined later	

# Instruction F with All Four Flags (so far)

---

```
IF OV = 1 THEN          /* Overflow */
    PC ← NEWPC; MODE ← 1 /* ABEND */
ELSE IF MP = 1 THEN     /* Memory violation */
    PC ← NEWPC; MODE ← 1 /* ABEND */
ELSE IF PI = 1 THEN     /* Invalid instruction */
    PC ← NEWPC; MODE ← 1 /* ABEND */
ELSE IF I/O = 1 THEN    /* I/O interrupt */
    OLDPC ← PC
    /* Perform additional functions to save process state here */
    PC ← NEWPC
    MODE ← 1
END IF
DECODER ← 00
```

# More on System Calls

---

The supervisor can use both user and privileged instructions; user programs can't.

- This means a user program can't execute even simple system calls like `open()` or `read()` by itself
- Communication with the supervisor is required
- This is the purpose of software-generated interrupts, or system calls

In a system call, the user program voluntarily relinquishes control to the supervisor.

- The supervisor then services the interrupt generated by the system call
- The interrupt handler code, before returning control to the user program, will generally copy any information the program requested into memory space accessible to it

# Next Time: Program State and System Calls

---